


STAT 339: Statistical Theory

Bayesian Hypothesis Testing

Anthony Scotina



Bayes Factors

A Slide about R

Suppose I am interested in estimating the **proportion**, p , of students at Simmons who have heard of R.

- A colleague claims that this proportion is **over 1/3**, and I want to test how plausible this claim is.



I think p is over 40% but I'm not very confident in that estimate. Suppose we use the following hierarchy:


- **Prior:** $p \sim \text{Beta}(8, 12)$
- **Data:**
 $Y \mid p \sim \text{Binomial}(n = 30, p)$,
where we observe $Y = 12$

Using the *Beta-Binomial conjugacy*, our **posterior** distribution is

$$p \mid Y \sim \text{Beta}(8 + 12 = \mathbf{20}, 12 + 30 - 12 = \mathbf{30})$$

Is the "over 1/3" claim plausible?

It's looking good! 😎




```
# 95% Credible Interval  
qbeta(c(0.025, 0.975), 20, 30)
```

```
## [1] 0.2699671 0.5376145
```

Posterior Probability

While the **visual** of the posterior PDF and the 95% **credible interval** help, we can calculate the **posterior probability** of p being over $1/3$:



```
1 - pbeta(1/3, 20, 30)
```

```
## [1] 0.8317821
```

Bayesian Hypothesis Testing

We can frame our analysis as two competing **hypotheses**:

- **Null Hypothesis**, $H_0 : p \leq 0.33$ (p is at most 33%)
- **Alternative Hypothesis**, $H_A : p > 0.33$

We just calculated the posterior probability of the *alternative hypothesis*:

$$P(H_A \mid Y = 12) = 0.832$$

- Therefore, the posterior probability of the *null hypothesis* is just the **complement**!

$$P(H_0 \mid Y = 12) = 1 - 0.832 = 0.168$$

Posterior and Prior Odds

We can put these two probabilities together to form the **posterior odds**:

$$\text{posterior odds} = \frac{P(H_A | Y = 12)}{P(H_0 | Y = 12)} = \frac{0.832}{0.168} = \boxed{4.952}$$

- Our *posterior* assessment is that p is 4.952 times more likely to be above $1/3$ than to be below $1/3$.

What about the odds **prior** to observing data?? Let's calculate **prior odds**.

- **Recall**: Our prior for p was $p \sim \text{Beta}(8, 12)$. Then $P(H_A) = P(p > 1/3)$:

```
1 - pbeta(1/3, 8, 12)
```

```
## [1] 0.7206633
```

$$\text{prior odds} = \frac{P(H_A)}{P(H_0)} = \frac{0.721}{1 - 0.721} = \boxed{2.584}$$

Bayes Factors (BF)

By comparing the *posterior odds* to the *prior odds*, we can obtain the **Bayes Factor**:

$$BF = \frac{\text{posterior odds}}{\text{prior odds}} = \frac{P(H_A | Y)/P(H_0 | Y)}{P(H_A)/P(H_0)}$$

- The Bayes Factor will provide some insight into how much our understanding of R familiarity at Simmons evolved after observing sample data.

```
# Bayes Factor  
4.952/2.584 # posterior_odds / prior_odds
```

```
## [1] 1.916409
```

Because the BF is a *ratio*, we should compare it to **1**.

Bayes Factor Scenarios

```
# Bayes Factor  
4.952/2.584 # posterior_odds / prior_odds
```

```
## [1] 1.916409
```

From Bayes Rules!:

1. $BF = 1$: The plausibility of H_A *didn't* change in light of the observed data.
2. $BF > 1$: The plausibility of H_A *increased* in light of the observed data.
 - The greater the Bayes Factor, the more convincing the evidence for H_A .
3. $BF < 1$: The plausibility of H_A *decreased* in light of the observed data.

While the **posterior probability** of H_A , $P(H_A \mid Y = 12) = 0.832$, is quite high, the $BF < 1$ established that the plausibility of my colleague's claim has *decreased* in light of the observed data.

🚫 BF Cut-Offs 🚫

While the statistics community generally advocates *against* using hypothesis testing to make **rigid conclusions**, there are tables out there that provide cut-off values for BFs.

For example, from [Jeffreys \(1998\)](#):

- $BF < 1$: Null hypothesis supported 😞
- $1 < BF < 10^{1/2}$: Evidence against H_0 "not worth more than a bare mention" 🔥
- $10^{1/2} < BF < 10^1$: Evidence against H_0 "substantial" 🔥 🔥
- $10^1 < BF < 10^{3/2}$: Evidence against H_0 "strong" 🔥 🔥 🔥
- $10^{3/2} < BF < 10^2$: Evidence against H_0 "very strong" 🔥 🔥 🔥 🔥
- $BF > 10^2$: Evidence against H_0 "decisive" 🔥 🔥 🔥 🔥 🔥

Don't actually use cut-offs, though! Science is more *nuanced*, and each hypothesis test and associated context should be considered individually. 🔥 🔥 🔥 🔥 🔥 🔥

Be careful with BFs!

Bayes Factors themselves aren't bad *per se*.

- But **null hypothesis significance testing** can be, and Bayes Factors are often used in this context!

Whether we use the *frequentist* or *Bayesian* perspective, we'll usually be able to reject a "no effect" null hypothesis by **gathering enough data**.

Example

Let's repeat the previous exercise with:

(i) $n = 1000$ and $Y = 400$

(ii) $n = 5$ and $Y = 2$

(In either case, the sample proportion is 0.4...)

Example

(Bayes Rules! Exercise 8.9)

For parameter p suppose you have a $\text{Beta}(1,0.8)$ prior model and a $\text{Beta}(4,3)$ posterior. You wish to test the null hypothesis that $p \leq 0.4$ versus the alternative that $p > 0.4$.

[Moose pic to fill space 📌]

Two-Sided Tests

Two-Sided Hypotheses

Let's revisit the R familiarity example from earlier. But suppose now that we're interested in testing *whether or not* 50% of the students at Simmons have heard of R.

- In other words, we're testing a **two-sided hypothesis**:

$$H_0 : p = 0.5 \quad \text{versus} \quad H_A : p \neq 0.5$$

We've seen this type of test under the **frequentist perspective**, but we hit a roadblock when trying to apply *Bayesian* techniques...

While the Bayesian perspective allows us to calculate the *probability* that H_A or H_0 is true, we can't really do this when p is **continuous**!



Bayes Factors (BF)

Recall that, by comparing the *posterior odds* to the *prior odds*, we can obtain the **Bayes Factor**:

$$BF = \frac{\text{posterior odds}}{\text{prior odds}} = \frac{P(H_A | Y)/P(H_0 | Y)}{P(H_A)/P(H_0)}$$

However, when we try to calculate the **posterior probability** that H_0 is true, $P(H_0 | Y = 12)$, we get zero!

$$P(p = 0.33 | Y = 12) = \int_{0.33}^{0.33} f(p | y = 12) dp = 0$$

- Therefore, the posterior odds are

$$\frac{P(H_A | Y)}{P(H_0 | Y)} = \frac{1}{0} = \text{🤪🤪🤪}$$

Using Credible Intervals

While we *clearly* can't divide by zero, we could use **credible intervals** to devise an approach to handling two-sided Bayesian hypothesis tests.

Recall that a **95% credible interval** for p is (0.234, 0.496).

```
# 95% Credible Interval  
qbeta(c(0.025, 0.975), 18, 32)
```

```
## [1] 0.2342242 0.4963872
```

- Using this, we have a *small* degree of evidence that the H_A is true.
 - The hypothesized value of p , 0.5, falls *just outside* the credible interval.

While one might argue that 0.5 isn't *substantially* outside of the credible interval, we should define what "substantial" means ahead of time!

Bayesian Testing, with a Buffer

Rather than testing

$$H_0 : p = 0.5 \quad \text{versus} \quad H_A : p \neq 0.5,$$

what if instead we tested

$$H_0 : p \in (0.45, 0.55) \text{ versus } H_A : p \notin (0.45, 0.55)$$

With this **buffer** in place, we can more rigorously claim *uncertainty* in the plausibility of H_A .

- The hypothesized range of p , $(0.45, 0.55)$, overlaps with the credible interval, $(0.234, 0.496)$.
- We could also calculate the Bayes Factor, since we'll no longer be dividing by zero!